

Detecting “Fake News” on Facebook

Hannah Eyre

Zane Zakraisek

December 6, 2017

The term “Fake News” gained popularity during the United States’ 2016 Presidential Election to describe a rapidly spreading phenomena of news articles deliberately spreading false information and hoaxes, often through attention grabbing headlines or headlines that resemble legitimate sources (Hunt, 2016). It became particularly notorious on social media sites and Facebook in particular, where the top 20 articles from hoax sites and hyperpartisan blogs garnered more user interaction between August 1st and election day on November 8th than the top 20 articles from a variety of established news sources such as *The New York Times*, *Washington Post*, *Business Insider*, and Fox News (Silverman, 2016).

FactCheck.org, part of the Annberg Public Policy Center at the University of Pennsylvania, breaks down how an individual can identify fake news into eight parts (Kiely & Robertson, 2016):

1. Consider whether the source is credible.
2. Read beyond the headline.
3. Check whether the author is credible or real.
4. Check whether the article is recent.
5. Check whether the article is a joke or satire.
6. Consider your own biases and how they affect your judgment.
7. Check supporting sources (if any) and make sure they abide by the same rules.
8. Ask experts or fact-checking sites.

Rapidly spreading fake news articles have a range of consequences, one instance culminating in a gunman attacking a Washington, DC pizza parlor over allegations of a satanic child sex abuse ring centered around John Podesta, Hillary Clinton’s 2016 campaign manager, in a conspiracy called “Pizzagate” (Haag & Salam, 2017). Various politicians and government agencies in the United States and internationally have voiced opinions on what qualifies as fake news and what do do about it, but no consensus has been reached. Facebook was initially reluctant to admit there was any problem with fake news on the website, however Facebook’s CEO Mark Zuckerberg has since released a statement describing how they plan to deal with fake news in the future, including renaming the phenomenon “false news” (Zuckerberg, 2016).

As an increasingly global and hotly contested issue, we would like to explore what responsibility Facebook has in regards to these eight points. We will discuss whether they have a responsibility to develop tools to detect fake news based off these guidelines and, if these tools exist, whether they should be used to remove content from the site. We will be using two particular examples, Pizzagate as mentioned above and conspiracies spread in the hours and days following the Las Vegas shooting. Pizzagate can be used as an example of what happens when fake news reaches a critical mass of interest online, having large followings on Reddit, 4chan, Breitbart and Infowars that actively encouraged “citizen investigation.” The Las Vegas shooting conspiracies are a group of examples of the rapid spread of fake news, particularly when real information is unavailable or unclear. In the early morning hours following the October 1, 2017 shooting, the website 4chan organized a successful effort to manipulate trending topics across the web and spread misinformation before real information was distributed.

When analyzing any issue, it’s important to define what framework you’re working in, along with what premises your arguments are based on. During our analysis of the Facebook’s role in moderating fake news, we’ll be working under a utilitarian framework. In particular, negative utilitarianism aimed at reducing the negative utility fake news causes. If the role that fake news has played in the events mentioned is examined, it is evident that the outcome results in a negative utility for Facebook’s users. News on Facebook has the potential to result in positive utility, and like fake news it is due to the rapid spread and wide range of people reached. We believe access to a variety of news sources is an important part of interacting with social media and results in a greater positive utility than fake news causes negative, so we believe removing news from Facebook would increase negative utility overall. As a result, we believe that it’s far more important to try to minimize the corresponding negative utility caused by fake news.

As we address each of these eight points mentioned above, the format shall be as follows. We’ll first examine the status quo where Facebook has not publicly addressed efforts made to combat fake news and how that is affecting users. Next, we’ll examine the outcome if Facebook chose to implement a fake news monitoring system. Lastly, we’ll weigh the different utilities and select the one that we feel minimizes the negative utility. In addition to this, we’ll also be examining whether or not the method is technically feasible to implement with current technology.

1 Credible sources: Sites and Authors

The first method in combating the spreading of fake news is to check whether the article in question comes from a credible source. This involves not just the author, the site itself. To start off, we’ll examine the utility if Facebook chose *not to* implement this source credibility check. Although Facebook currently implements a very light form of censorship (citep), the scenario would be quite similar to how Facebook currently monitors articles, which is very little. While most articles that

Facebook feeds users are based off of topics that they've shown interest in, nearly every Facebook user has probably been shown an article from an underground website that someone on their friends list has liked or shared. If Facebook didn't implement source credibility prioritization, then the effect of one person sharing an article like this could have quite a negative outcome. The article in question can easily make its way to hundreds of people. For most fake news articles, this has the effect of increasing the amount of negative utility, as most fake news articles are spread with negative motives.

Now assume Facebook was to implement source credibility prioritization. In this same scenario, if someone was to like or share a post, Facebook could first check its feedback for that particular site and/or author. If there is enough feedback, and the score is negative, Facebook could deprioritize the article on other people's news feeds by a factor of the magnitude of the negative score. Likewise, if the site or author has a positive score, Facebook could choose to display the article as it normally would, or maybe even prioritize it higher. From a utility perspective, this has the potential to significantly reduce negative utility brought about by fake citemps. On the other hand, there does exist the potential to filter out some citemps which many people may find to be *good*. In other words, this has the potential to significantly decrease negative utility, with a smaller side effect of decreasing some positive utility. From an overall perspective however, implementing source credibility prioritization has the net effect of decreasing negative utility

In terms of the technical feasibility of this point, we feel like it one of the lowest hanging fruit in the pursuit of diminishing fake news. Once possibility would involve Facebook storing a few metrics for a site based upon feedback received by readers. Facebook already has a way to flag particular articles as inappropriate. In the same way, Facebook could allow a reader to optionally rate the credibility of the article or the site in general, and use these ratings to prioritize and deprioritize sites.

Lastly, there's the idea of checking whether or not the author of an article is indeed a real person, and is indeed who they claim to be. We feel that this would be a much more difficult task for Facebook to implement. Additionally, while author authentication is a concept that would indeed help lower negative utility, the implementation of it would most likely cross the privacy line. This in turn increases negative utility. For example, Facebook may attempt to link the author of an article on some site to an individual account on Facebook for an added level of author verification. Since this has implications in terms of privacy, we feel that author authentication is not an angle that Facebook should approach this problem from. (hannah- what about the "verified user" thing that already exists on facebook, could that be used? make verification for that checkmark more stringent or up to some determined "fake news standards", then only allow topics from verified accounts onto trending? has problems with viral user posts though)

2 Reading Beyond the Headline

When posting an external link to Facebook, information is drawn from the content inside meta elements of the page source by Facebook's web crawler and placed into a preview on the post itself. If an element labeled according to the Open Graph protocol, the crawler will place it in the link preview on the site (Facebook, n.d.). With this system in place, every linked website will appear the same on a users news feed. Any site can determine how its posts show up when linked on Facebook, regardless of the quality of the content or site itself.

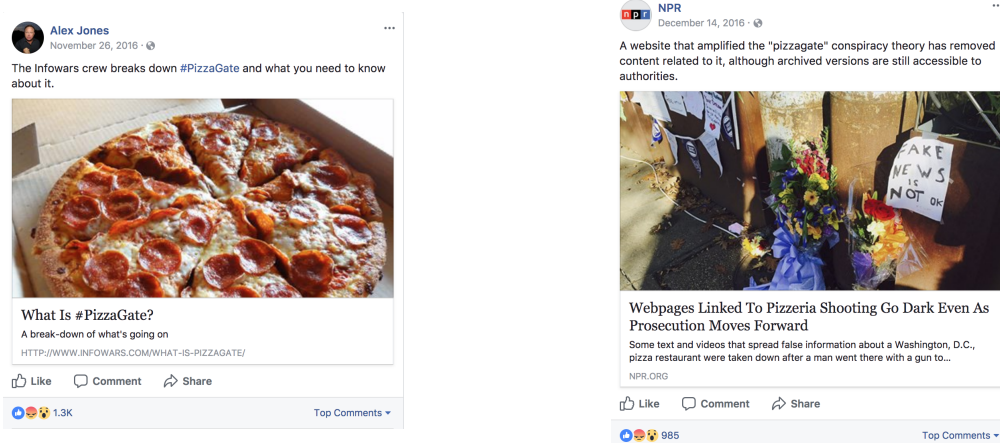


Figure 1: A comparison of posts by Alex Jones, creator of *Infowars*, (Jones, 2016) and *NPR* (NPR, 2016) about Pizzagate.

This uniform formatting treats *NPR* and *Infowars* or any other site that uses the meta elements according to Facebook's specifications as equals. It places less emphasis on the source of the article than on the headline and picture. The examples above are posted from official accounts, allowing users to easily see where the article comes from. However, the same articles posted by friends or unofficial pages make the source more difficult to determine. In addition, 50% of shared links generate less than 1% of clicks on social media (Gabiello, Ramachandran, Chaintreau, & Legout, 2016), meaning most shares happen without the user clicking the link and reading the article. Therefore what a user sees on Facebook's preview is most likely all that user sees and is entirely determined by the web developer, with the possibility that the headline, picture, and description are totally unrelated to the actual content of the article, factually incorrect, or telling an attention-grabbing but incomplete story.

The ability to spread misinformation through headlines, pictures, and short descriptions while knowing that users rarely click on articles to reach

Facebook could resolve this by implementing an automatic summarization of the contents of the page, however summarizing text has many technical hurdles

3 Age of the Article: Hannah

Considering the age of the article is an important aspect of identifying fake news. Both very recent articles and older articles can be subject to scrutiny. Older articles can be brought back lacking any date other than the age of the post using Facebook's link previewing feature or quoted by newer articles out of context with no means of cross referencing, manipulating a potential reader into thinking an article is relevant to a current situation in addition to whether they are factually accurate. Likewise, news articles that rapidly accumulate reactions, shares, and likes can appear on an unmoderated trending topics list, regardless of the content of the article. On October 1, 2017, a man opened fire on a crowd attending a concert in Las Vegas, killing 59 people and injuring over 500. Almost instantly, fake news began to spread about the topic, reaching the top of trending pages across the internet, including Facebook. The articles and posts varied in topics, ranging from flooding social media with posts falsely searching for missing relatives, calling the shooter an Islamic convert and the shooting an act of Islamic terror, or saying the shooter was involved with Antifa and had attended anti-Trump rallies (Boderick, 2017). All of these false statements were posted and promoted in a way that was meant to disrupt the flow of accurate information or make those seeking help less able to find it.

4 Article Genre: Hannah

the onion, dedicated satire site, <https://www.theonion.com/no-way-to-prevent-this-says-only-nation-where-this-r-1819580358>

new yorker, sometime satire,

"just a joke" defense for harassment,

5 Interpretation Bias: Hannah

Facebook recently came under fire for their manual curation of trending news topics when former Facebook employees were interviewed by Gizmodo about the trending news section. One curator said "There was no real standard for measuring what qualified as news and what didn't" and that they regularly avoided right wing sites such as *Breitbart* and *The Blaze*. When pressed for a comment, Facebook responded by claiming they "take allegations of bias very seriously" (Nunez, 2016). By framing bias as a negative, Facebook implies that they view themselves as a neutral entity and that they believe it is possible for the website and trending topics to be free of bias.

trending topic staff fired, <https://www.theguardian.com/technology/2016/aug/29/facebook-fires-trending-topics-team-algorithm>

6 Cross Reference: Zane

7 Supporting sources: Zane

The final method for identifying fake news is to check supporting sources (if any) and make sure they abide by the same rules. In other words, this method attacks the problem with a recursive approach. When a user shares an article on the site, the fake news algorithm could take a depth first approach at calculating a trustworthiness metric. This can be done by first calculating the trustworthiest of the sources linked to in the article, and then using this metric in the overall trustworthiness calculation for the main article.

From a utilitarian perspective where our main goal is to minimize the amount of negative utility brought about by fake articles, we feel that implementing this method has a few angles to consider. Since this method is basically taking the other methods that we feel are appropriate, and applying them in a recursive manner, we feel that for the most part this would be an extra step in minimizing the spread of fake news. One point to keep in mind however is that by taking a recursive approach, we may be rating the credibility of one page based on the credibility of a page with multiple degrees of separation that have virtually nothing to do with each other. Also, as mentioned in *Credible citeps and Sources*, we don't have a way to flag a link with the context that it's being cited in. For example, if an author is citing a fake news article with the intention of pointing out the fake news, the trustworthiness metric should somehow understand this context, and not use this against the articles score.

In terms of the technical feasibility of the implementation, there are a few issues that would make this point difficult to implement. The first hurdle is that there would have to be some kind of standardization to the method that sites use to track citations and references. Most sites that people read have a nonstandard citations section at the bottom of the page. Although this conveys the required information to the *reader*, it fails to be parsable by a web crawler. Additionally, this depth first approach to calculating source credibility would involve storing a large amount of metadata. LINK SHORTENERS (hannah - might want to talk about sentiment analysis for the linking fake news part)

8 Conclusion

With the increasing veracity of fake news across all of social media, particularly Facebook, and the consequences of letting misinformation spread rapidly and unchallenged, Facebook has a responsibility to come up with an effective and equitable strategy to moderate it. Facebook has an enormous platform, a Pew Research Center survey found that 66% of respondents said they use Facebook, 68% of those users get news from Facebook, but only 33% of Facebook users get news

from local TV or dedicated news sites and apps, the next most popular categories after Facebook itself (Shearer & Gottfried, 2017). Facebook earns more than twice the engagement in news than any single non-social media platform but Facebook does not hold itself to any standard for the integrity of their sources unlike most traditional formats. (hannah - maybe on the moderation algorithm part? like lots of negatives to developing one but is fake news by itself worse? what conclusion should we draw here. Tell facebook to spend more time investigating?)

There are many challenges involved in implementing a method of fake news moderation. With the volume of posts posted and shared, it is unlikely Facebook will be able for human employees to keep up when moderating content. Algorithmic intervention seems to be the most practical solution if human moderation and automatic moderation were equally accurate, however there are just as many different struggles for developing an algorithmic fake news detection system. **...describe problems here, privacy/false positives/false negatives, etc...**

If Facebook were to decide to invest in an automated system for detecting fake news, they will need to be transparent at every step of the process. In order to begin designing such a system, Facebook needs to determine what they value in a news source and be willing to stand by those beliefs in the face of adversity or openly change them in response to criticism. Their negative response towards “allegations of bias” in response to not allowing a site that frequently promotes white nationalism, fake news, and harassment campaigns, *Brietbart*, onto the trending topics list, they are themselves making a political statement in tacit support of the platform’s message by treating its inadvertent suppression as a something to be corrected, which conflicts with previous statements against hate speech and harassment. Neutrality is not apolitical. **more stuff**

References

- Boderick, R. (2017, October 2). Here are all the hoaxes being spread about the las vegas shooting. *Buzzfeed News*. Retrieved from <https://www.buzzfeed.com/ryanhatesthis/here-are-all-the-hoaxes-being-spread-about-the-las-vegas>
- Facebook. (n.d.). A guide to sharing for webmasters. Retrieved from <https://developers.facebook.com/docs/sharing/webmasters/>
- Gabielkov, M., Ramachandran, A., Chaintreau, A., & Legout, A. (2016, June). Social clicks: What and who gets read on twitter? *ACM SIGMETRICS / IFIP Performance 2016*.
- Haag, M. & Salam, M. (2017, June 22). Gunman in ‘pizzagate’ shooting is sentenced to 4 years in prison. *The New York Times*.
- Hunt, E. (2016, December 17). What is fake news? how to spot it and what you can do to stop it. *The Guardian*.
- Jones, A. (2016, November 26). Retrieved from <https://www.facebook.com/AlexanderEmerickJones/posts/10154777113688459>
- Kiely, E. & Robertson, L. (2016, November 16). How to spot fake news. *FactCheck.org*. Retrieved from <http://www.factcheck.org/2016/11/how-to-spot-fake-news/>
- NPR. (2016, December 14). Retrieved from <https://www.facebook.com/NPR/posts/10155078014821756>
- Nunez, M. (2016, May 3). Want to know what facebook really thinks of journalists? here’s what happened when it hired some. *Gizmodo*. Retrieved from <https://gizmodo.com/want-to-know-what-facebook-really-thinks-of-journalists-1773916117>
- Shearer, E. & Gottfried, J. (2017, September 7). News use across social media platforms 2017. *Pew Research Center*. Retrieved from <http://www.journalism.org/2017/09/07/news-use-across-social-media-platforms-2017/>
- Silverman, C. (2016, November 16). This analysis shows how viral fake election news stories outperformed real news on facebook. *BuzzFeed News*. Retrieved from <https://www.buzzfeed.com/craigsilverman/viral-fake-election-news-outperformed-real-news-on-facebook>
- Zuckerberg, M. (2016, November 18). Retrieved from <https://www.facebook.com/zuck/posts/10103269806149061>